

Specify your system

One page from “we tested it” to “here is the property”

1 • The decision your AI makes that, if wrong, hurts someone

2 • Its inputs, with units and valid ranges (this is the *region*)

3 • Its outputs (classes, scores, actions)

4 • One property, in English

“For every _____ such that _____, the system must _____.”

5 • The same property, as a specification (template)

template

```
@network
f : Input -> Output

valid : Input -> Bool
valid x = ____ ≤ x ! 0 ≤ ____ and ...

@property
p : Bool
p = forall x . valid x and _____
    => _____
```

6 • What would a counterexample look like, and who must see it?

7 • What cannot be verified this way (meaning, scale, adversaries, agents) and what you will do instead (testing, monitoring, human oversight)

8 • Who decides the thresholds and the margin, and who signs?

Bring this sheet to your Beta demo: it reads as an assurance argument. The verification tools for a small tabular or control network are in the notebooks; for anything else, the sheet is still the right first step.